Binary search trees (BST)

at every node, all nodes to the left have smaller key values

right larger

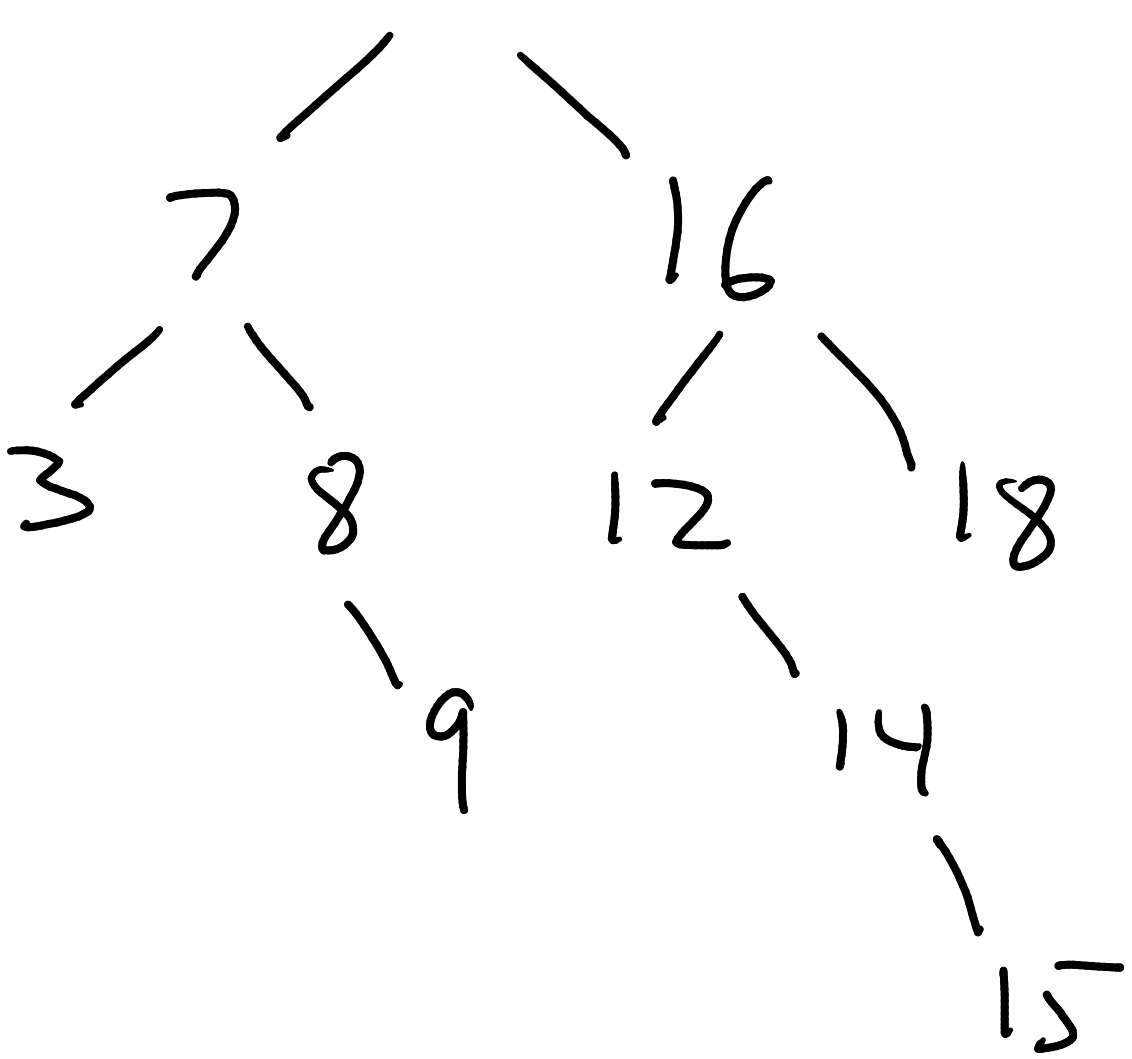

Is this a correct BST?

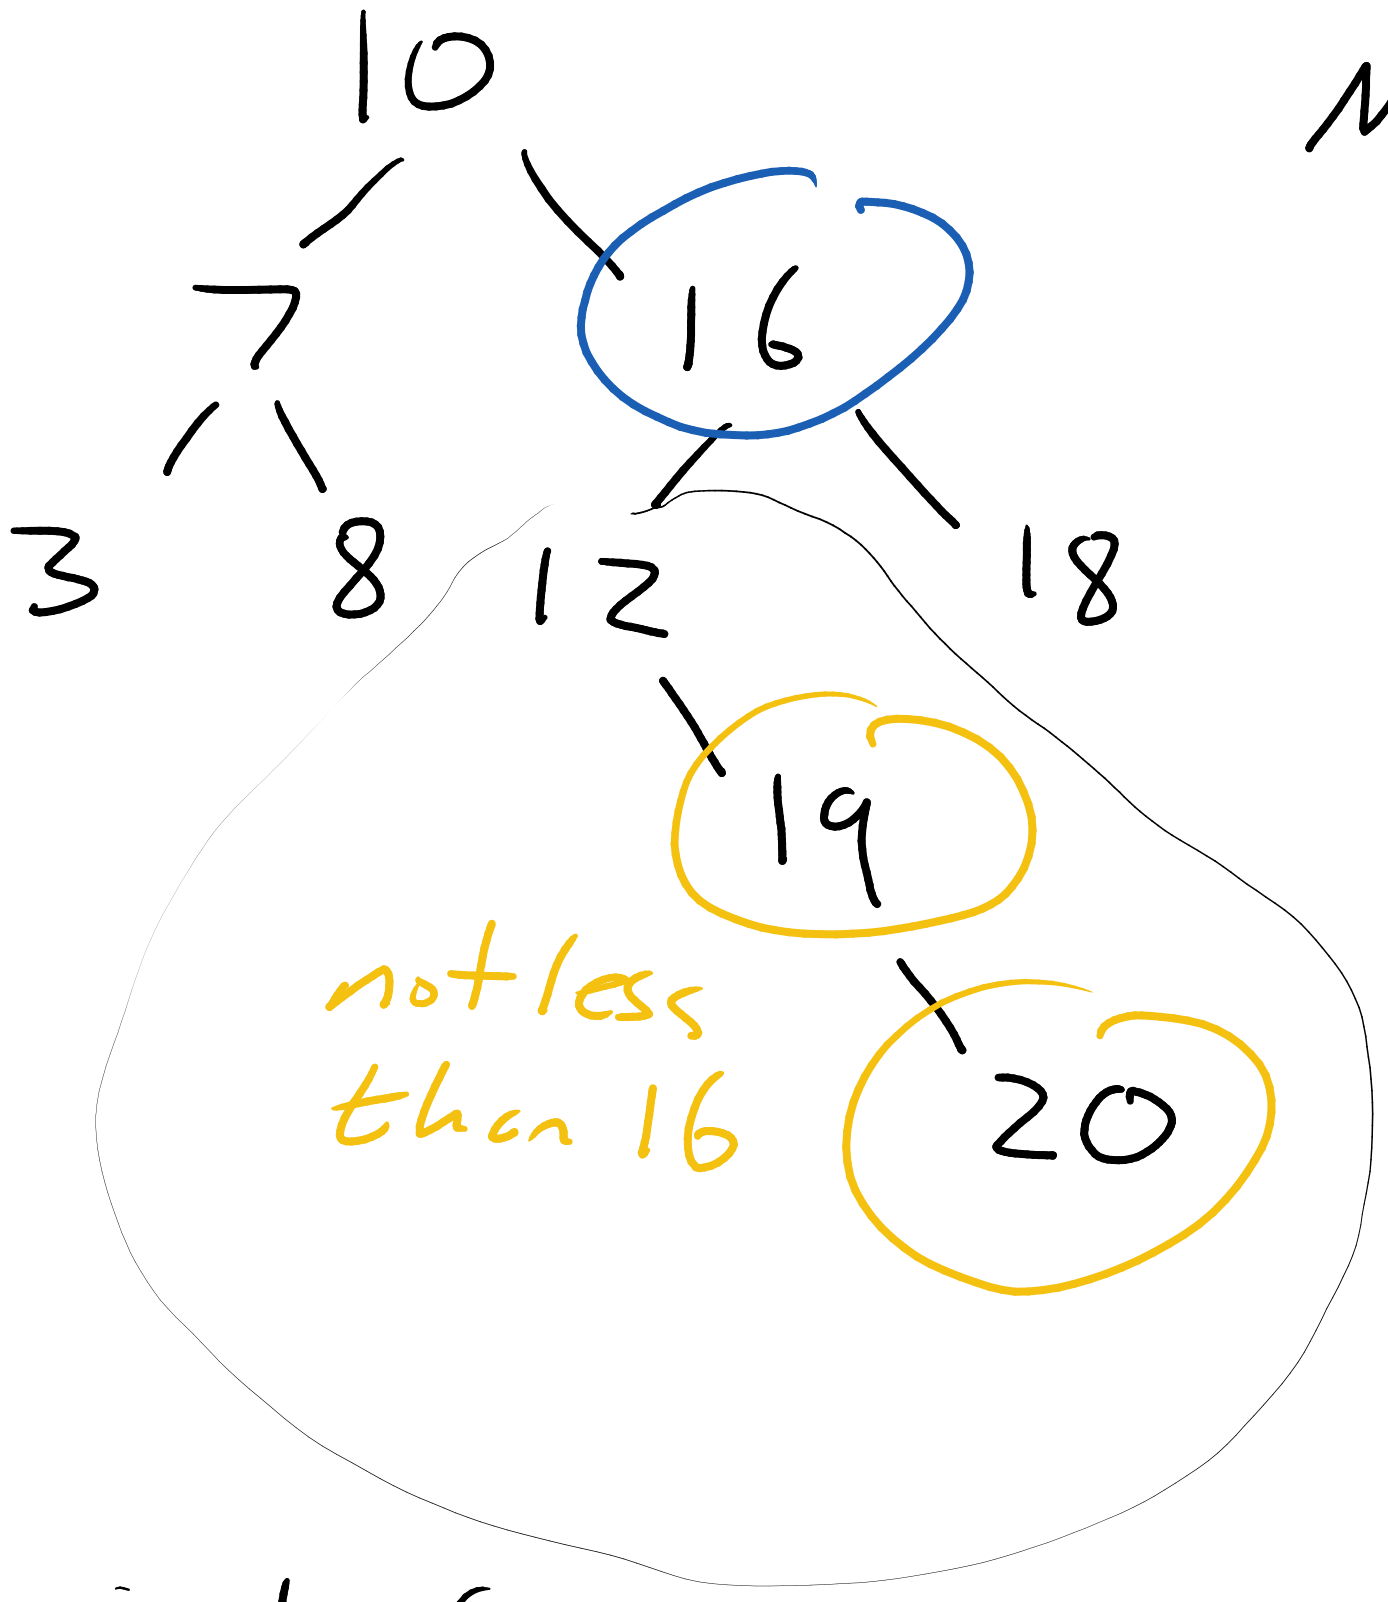

No!

Is it just for numbes?
No, it's for anything you can order

Strings

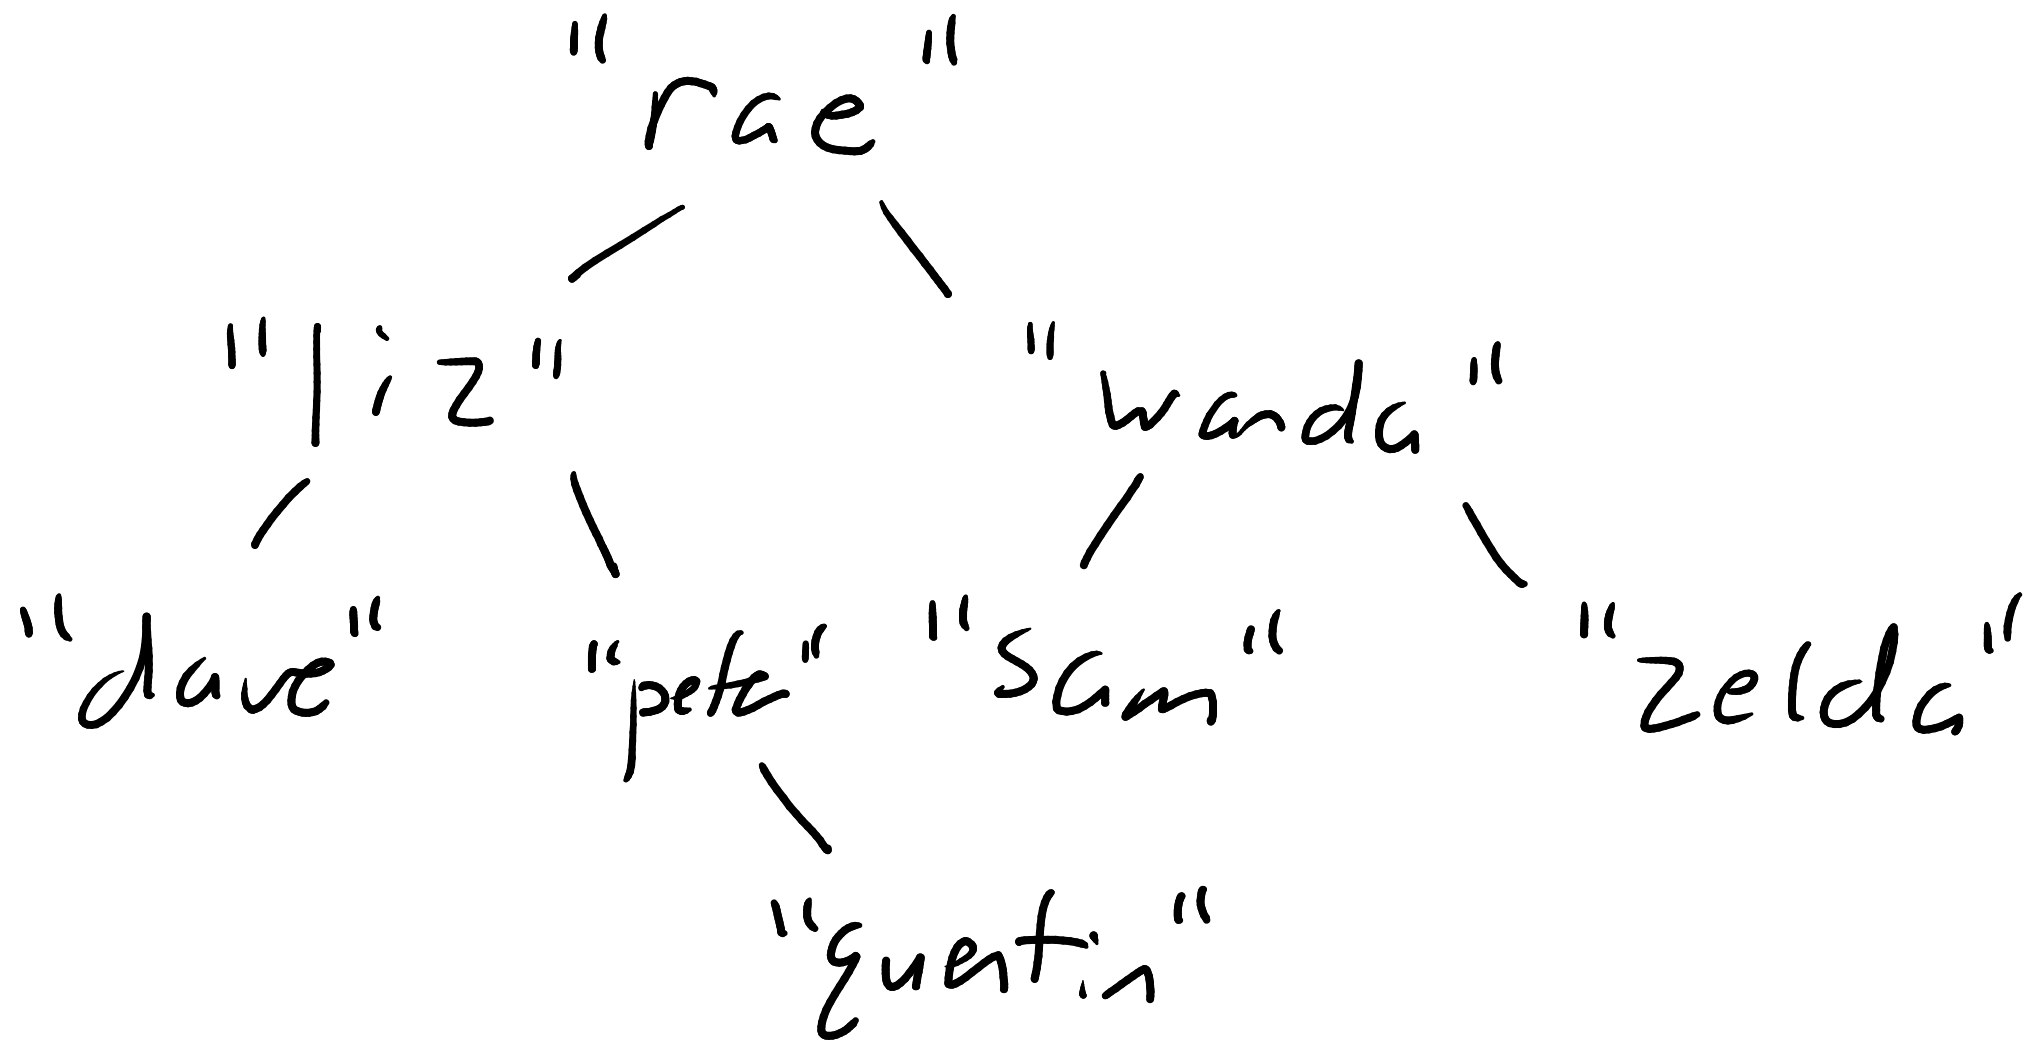

Why? Useful for sets and maps, just like hash tables (dicts)

For a map, just store a value next to the key

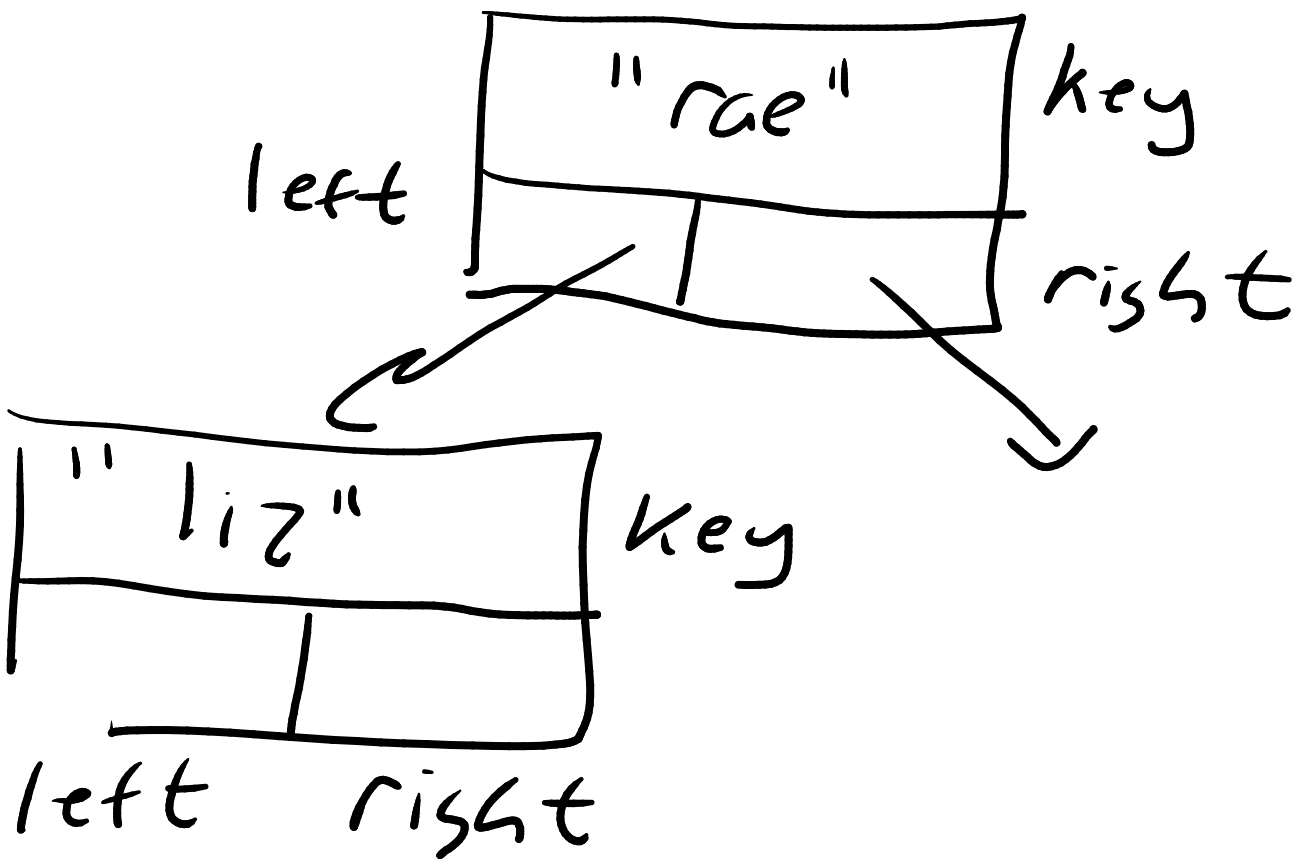

For a map, just also store a value

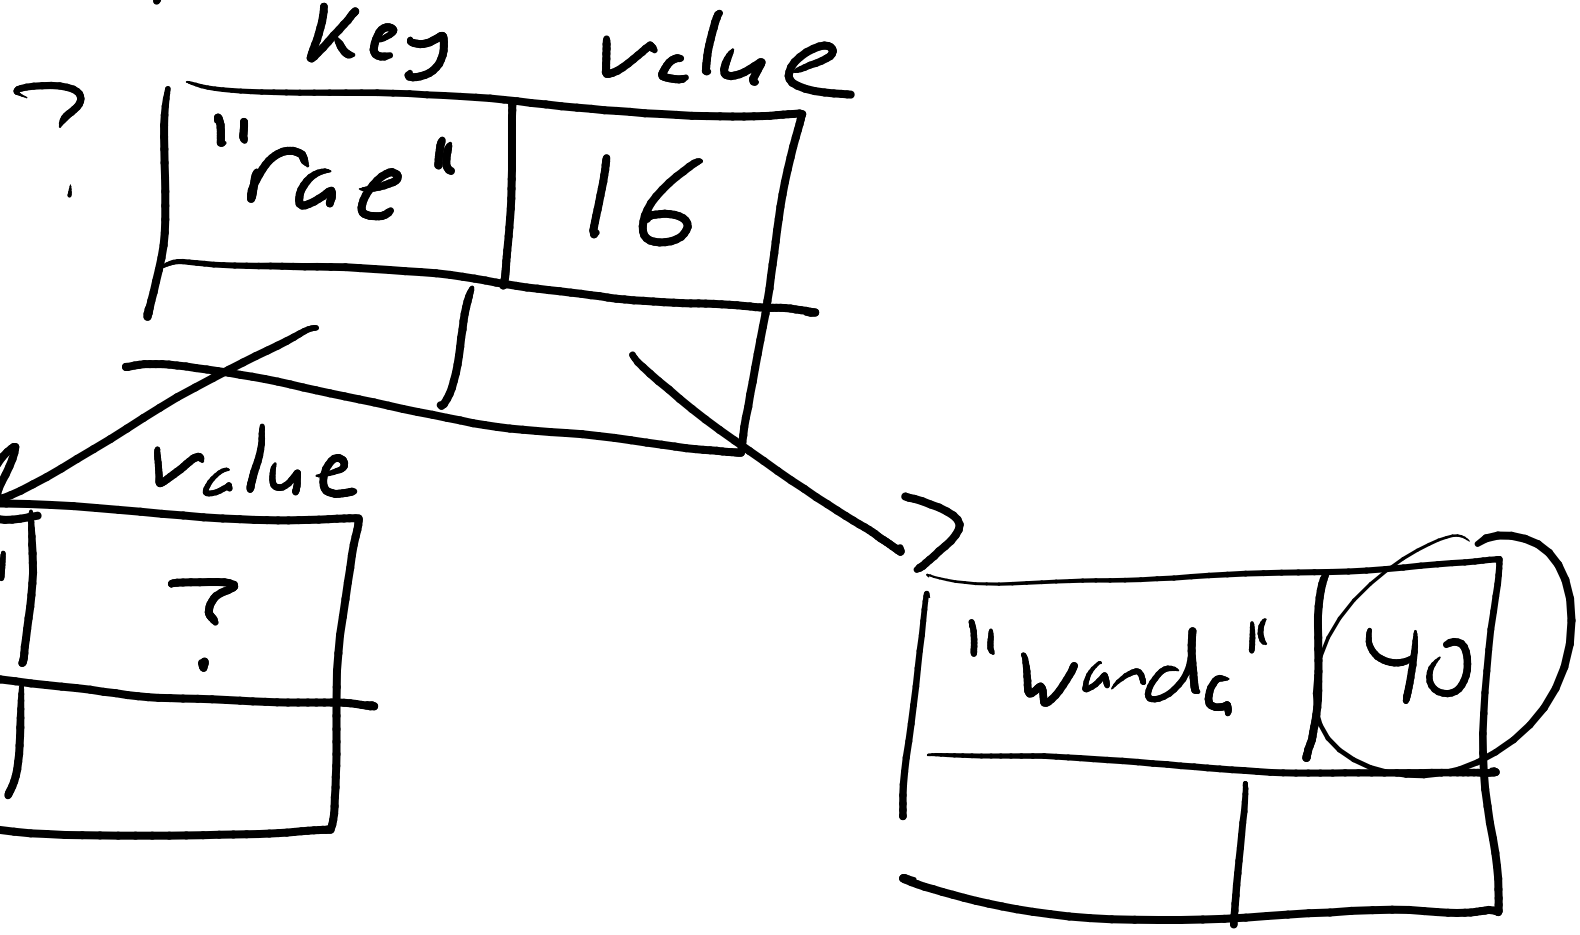

Lookup "wanda"

Performance

For n items, what is worst-case number of ops for insert/delete/lookup

- all 3 of these potatically involve going out to a leaf

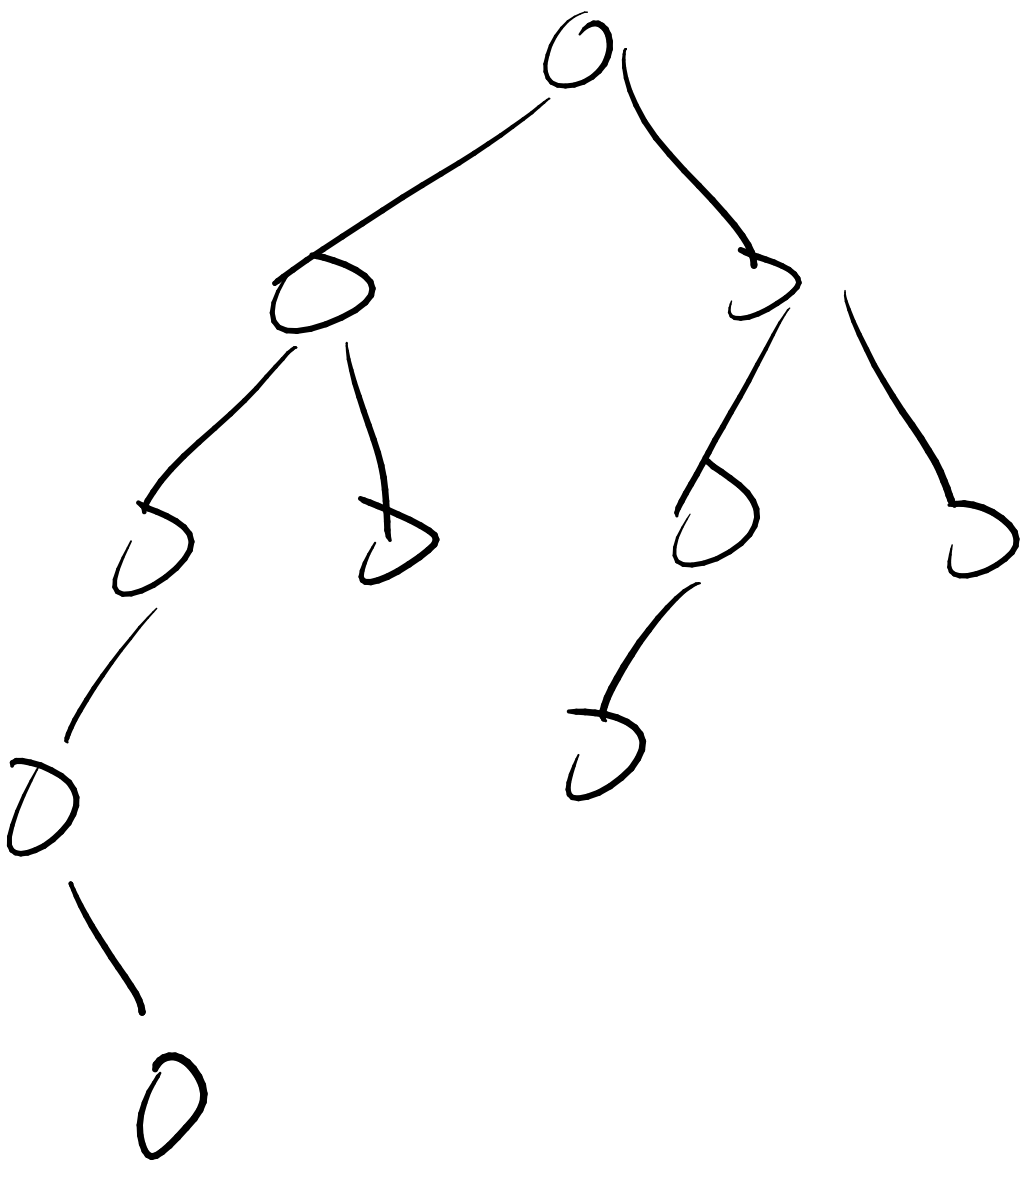

Worst case, how deep is tree?

Two cases:

- if tree is balanced
- if tree is not balanced

Balanced: all leaves on same [or off-by-one] level

(oversimplified defn, not quite right)

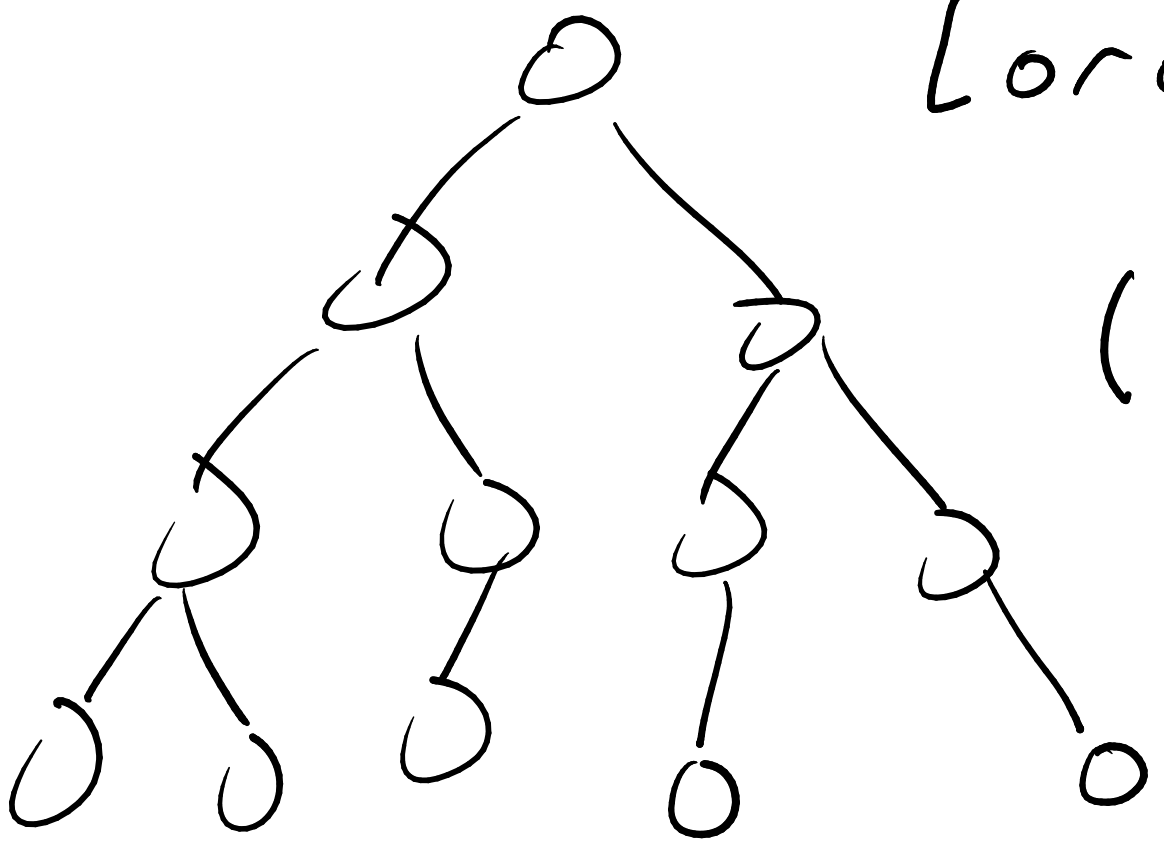

If I have n items,

approx how deep is the tree if balanced?

(how many levels?)

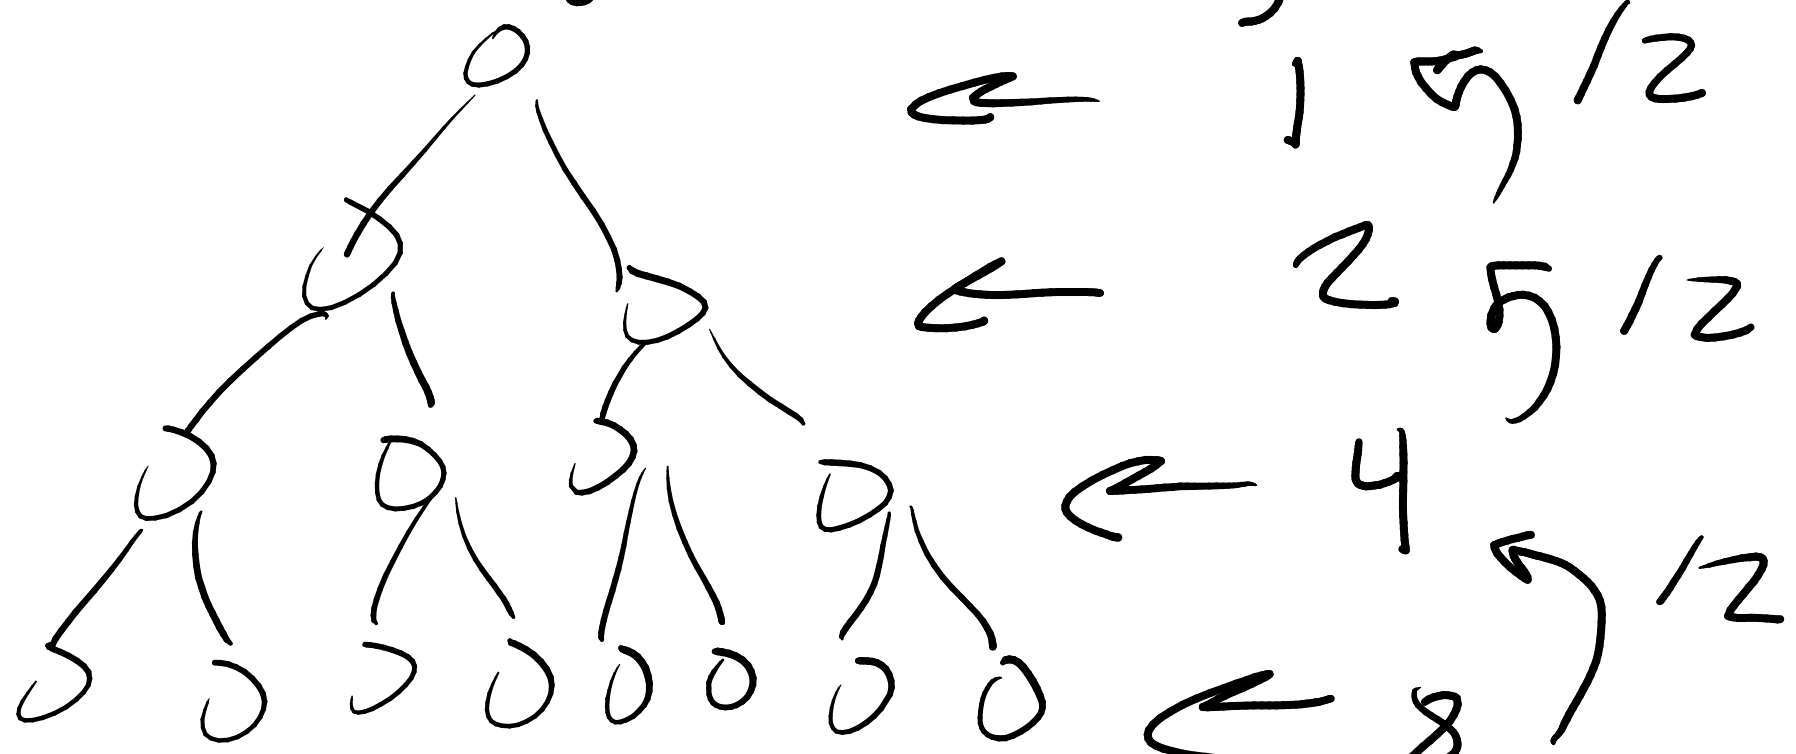

bottom row is approx half the total

15

Number of levels is the number of times it can keep cutting n in half (approx) until I get 1

So in general, the height of tree is approx $\log_2 n$

Time to do insert/delete/lookup is $O(\log n)$

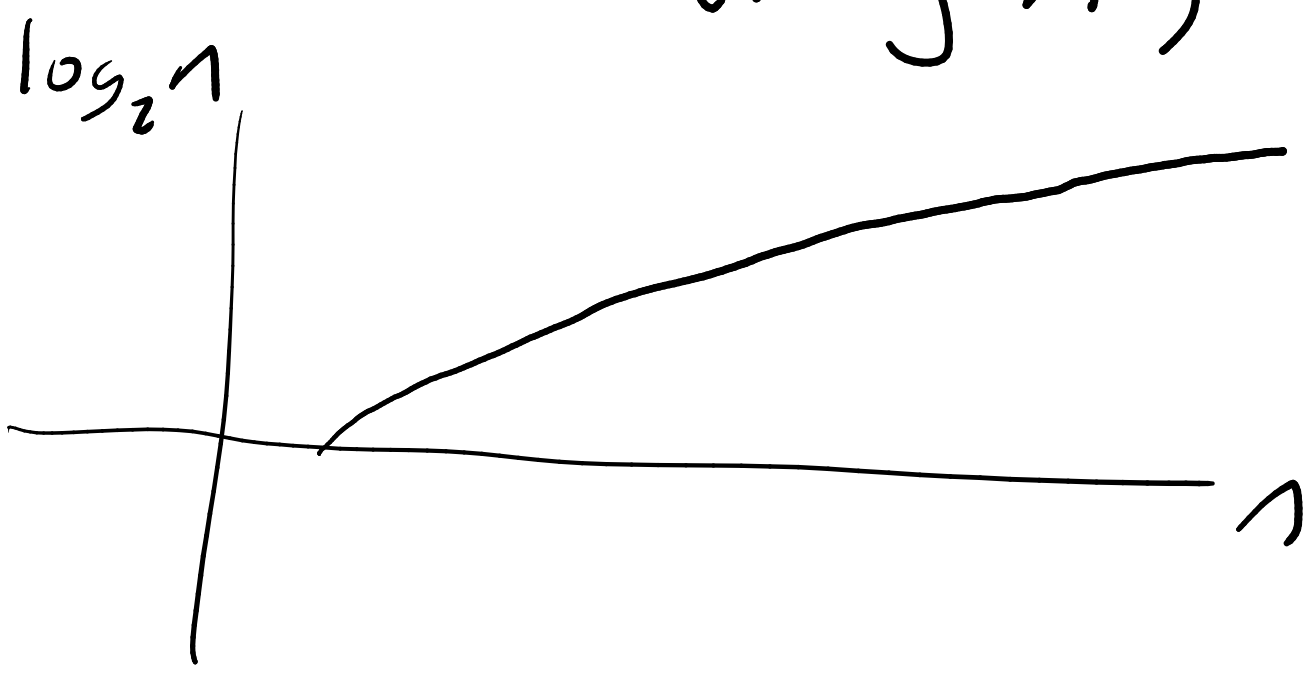

If unbalanced, worst case is a disaster.

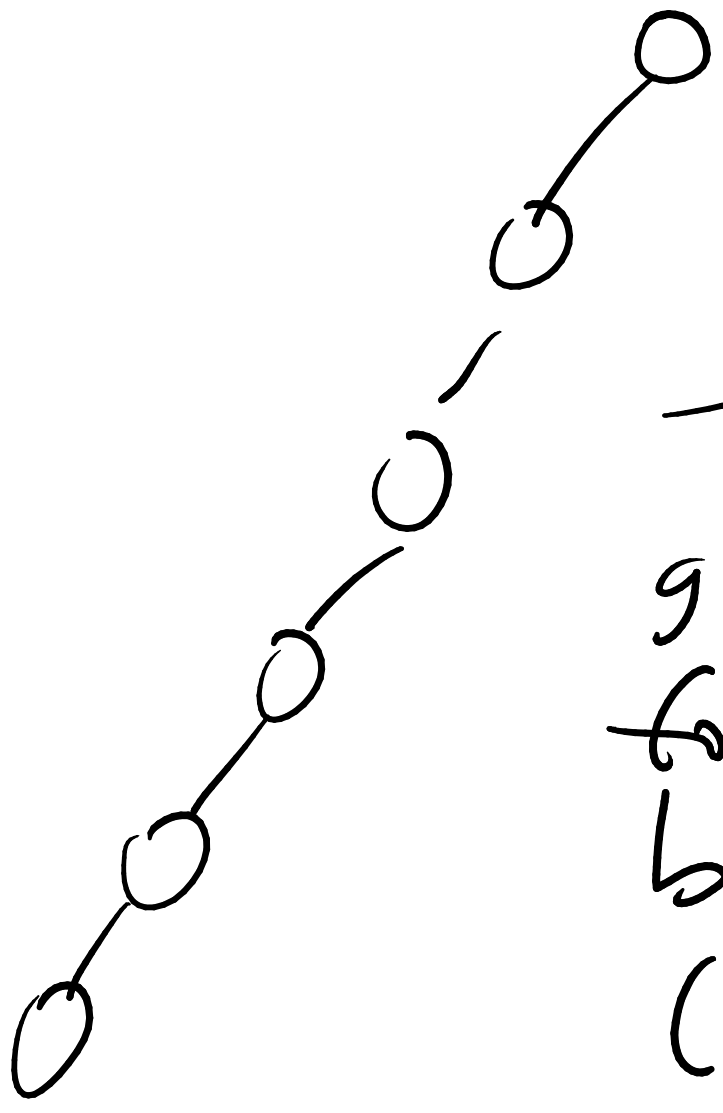

Worst case time is $O(n)$

There are lots of great approaches for forcing a tree to be balanced. [AVL trees] (not today)

In practice, you can get $O(\log n)$ performance.

How does this compare w/ hashing?

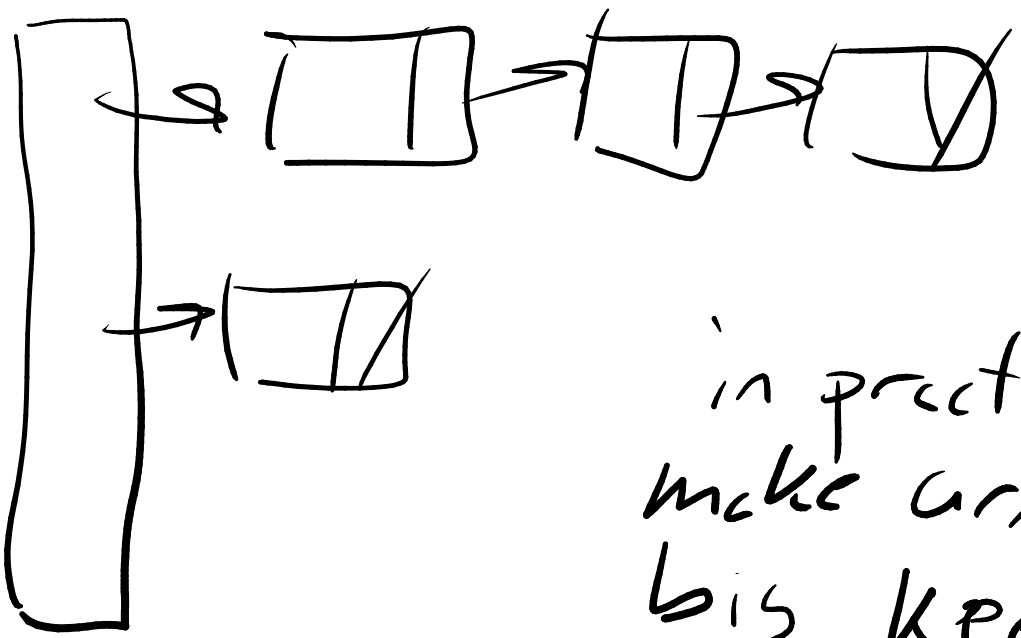

In practice

$O(\lambda)$

$\lambda \leq \frac{\#\text{ of items}}{\text{size of array}}$

in practice, make array big, keep λ small, $O(1)$

Hashing: $O(1)$ ⇒ in practice,
2, or 3, always
(just make array big enough)

BSTs, in practice
$O(\log n)$

Hashing is definitely faster
(when you do it right)
- you need to get the size of the array right

Real main advantage of BST:
"Get me all keys > 30"
values for

In a BST look for 30:

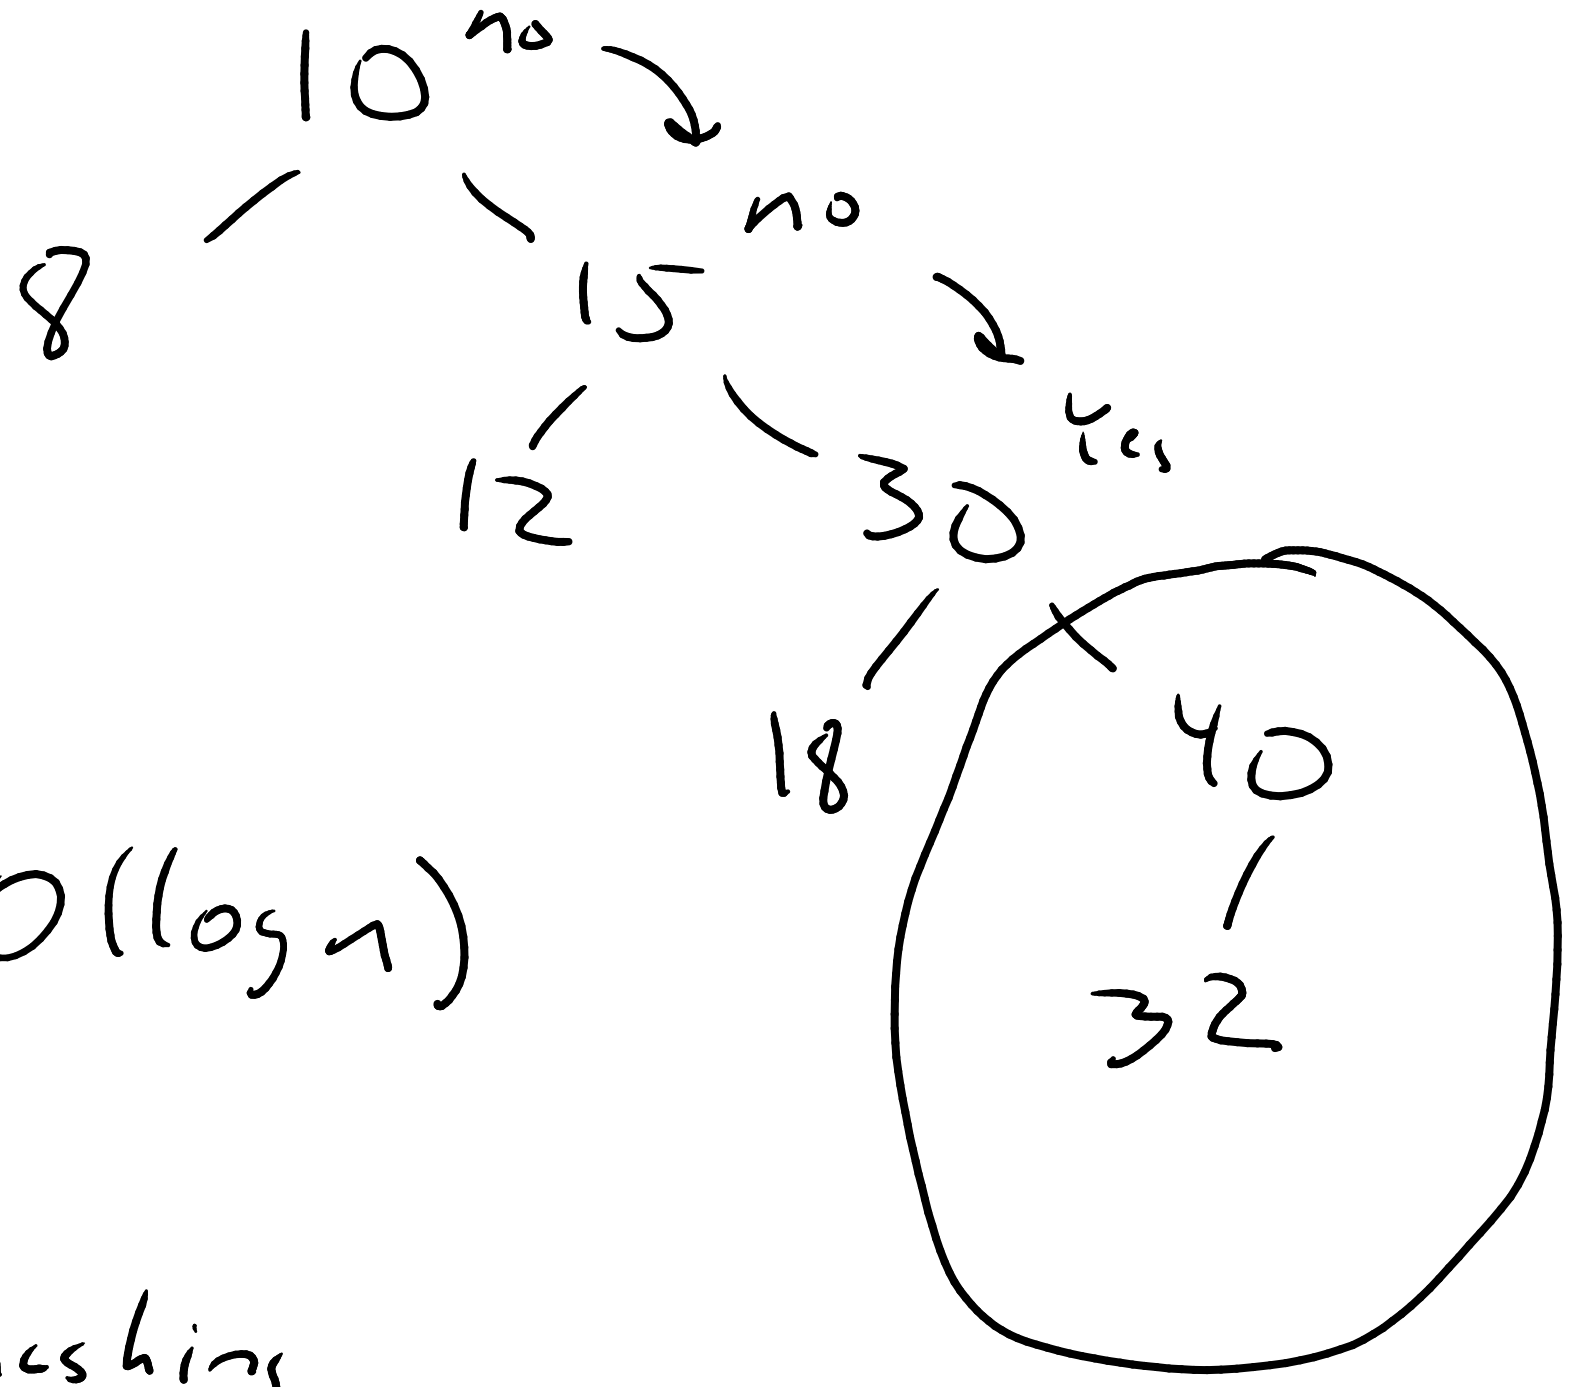

O(log n)

with hashing

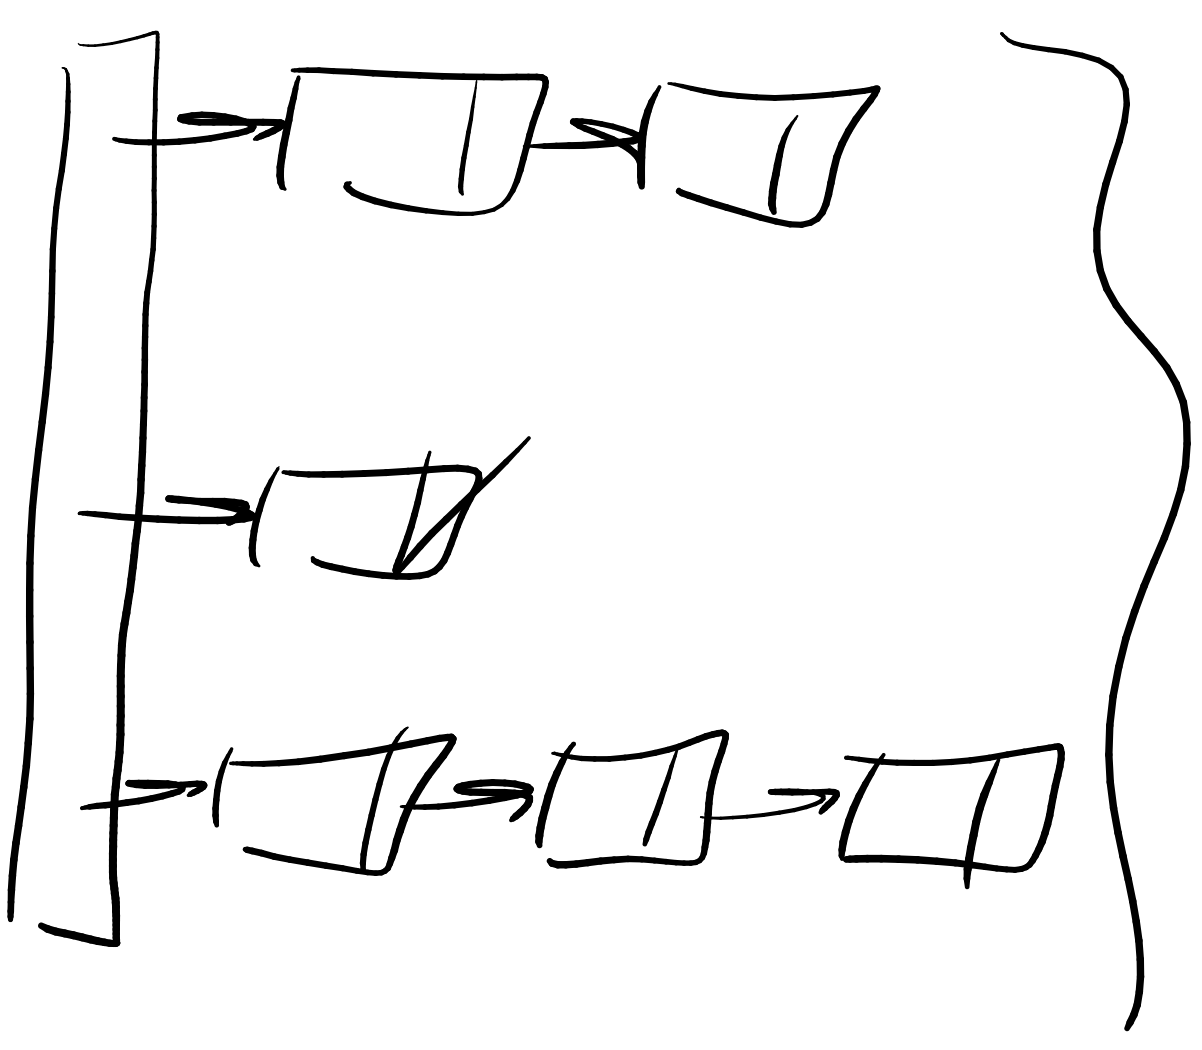

have to look at everything

O(n)